

COMP4901K/Math4824B

Machine Learning for Natural Language Processing

Lecture 1: Introduction

Instructor: Yangqiu Song

Logistics

- Instructor: Yangqiu Song
- Email: yqsong@cse.ust.hk
- Canvas (<https://canvas.ust.hk>)
 - Lecture notes
 - Assignments
 - Projects
- TAs
 - Tianqing Fang: tfangaa@connect.ust.hk
 - Zizheng Lin: zlinai@connect.ust.hk

Background of this Course

- Purpose of this course
 - Currently a topic course for senior undergraduate students
 - Enrich their experience on text data analytics
 - Equip them with powerful analytic tools for future career
 - Target to be a future elective course for the Data Science and Technology (DSCT) program
 - <http://dsct.ust.hk/>
- Shared Course Codes COMP4901K/Math4824B
 - Math students can replace final project with a survey paper

More about this Course

- The course covers the knowledge from both CSE and Math areas:
 - The fundamental machine learning models that deal with real natural language processing problems.
 - The programming and analytic tools such as Python to deal with real natural language processing problems.
 - Many example code will be provided
 - The problem solving skill for real natural language processing problems, such as building end to end system to deal with a certain problem.
- So
 - It's not a pure machine learning course
 - It's not a pure NLP course
 - We can cover more modern and advanced learning algorithms for NLP tasks
 - Will be under “*Restricted Deep Learning Electives*” in CSE programs
 - Among the courses needed to satisfy our Area requirement (for a total of min. 15 credits), only one course among these “Restricted DL Electives” can be used to count as one of those 5 courses.
- Prerequisites
 - Linear Algebra (MATH 2111 or MATH 2121)
 - Programming (COMP 2011/2012)
 - Discrete Math (COMP 2711 or Math 2343)

More about this Course

- It's a data science course
- We will have a lot of practicing opportunities to make sense of using text data
 - To understand text data problems arising in the areas of commerce and industry etc.
 - To model text data problems using different mathematical tools.
 - To design and implement efficient algorithms to solve different NLP problems.
 - To interpret the results provided by different algorithms and apply them to the data problems to gain meaningful insights or offer predictions.

Goals of this Course


- Learn about the problems and possibilities of natural language analysis:
 - What are the major issues?
 - What are the major solutions?
 - How well do they work?
 - How do they work?
- At the end you should:
 - Agree that language is subtle and interesting!
 - Feel some ownership over the algorithms
 - Be able to assess NLP problems
 - Know which solutions to apply when, and how
 - Be able to read papers in the field

Today

- Why NLP is difficult
- How to solve it (in general)?
- What we'll do in this course


Natural Language

- Understanding language is a very complex thing
- But something that humans are amazingly good at



Artificial Intelligence: Turing Test

- Replacement of “Can machines think?”
 - Behavioral test
 - not just natural language understanding
- Natural language conversation
 - *Imitation game (1950)*
 - Human?
 - Computer?




The AI Winter

- AI winter: 1974–80 and 1987–93
 - 1966: the failure of **machine translation**,
 - 1970: the abandonment of **connectionism**,
 - 1971–75: DARPA's frustration with the **Speech Understanding Research** program at Carnegie Mellon University,
 - 1973: the large decrease in AI research in the United Kingdom in response to the Lighthill report,
 - 1973–74: DARPA's cutbacks to academic AI research in general,
 - 1987: the collapse of the Lisp machine market,
 - 1988: the cancellation of new spending on AI by the Strategic Computing Initiative,
 - 1993: expert systems slowly reaching the bottom, and
 - 1990s: the quiet disappearance of the fifth-generation computer project's original goals.

Enabled by Big Data

- 1971–75: DARPA's frustration with the **Speech Understanding**
- “Watson is a **question answering** (QA) computing system that IBM built to apply advanced
 - natural language processing,
 - information retrieval,
 - knowledge representation,
 - automated reasoning, and
 - machine learning technologies
- to the field of **open domain question answering**.”




In 2011, Watson competed on Jeopardy! against former winners Brad Rutter and Ken Jennings. Watson received the first place prize of \$1 million.




Enabled by Big Data



Startup Companies (2015)

Which Artificial Intelligence Categories Are Seeing the Most Innovation? by  Venture Scanner




Number of Exits (Acquisitions and IPOs, 2017)



ARTIFICIAL INTELLIGENCE
Exit Activity by Category

VS / VENTURE
SCANNER




Data as of July 2017


What's Special about Human Language?

- A human language is a discrete/symbolic/categorical signaling system
 - With very minor exceptions for expressive signaling (“I loooove it.” “Whoomppaaa”)
- Large vocabulary, symbolic encoding of words creates a problem for machine learning – sparsity!

Face

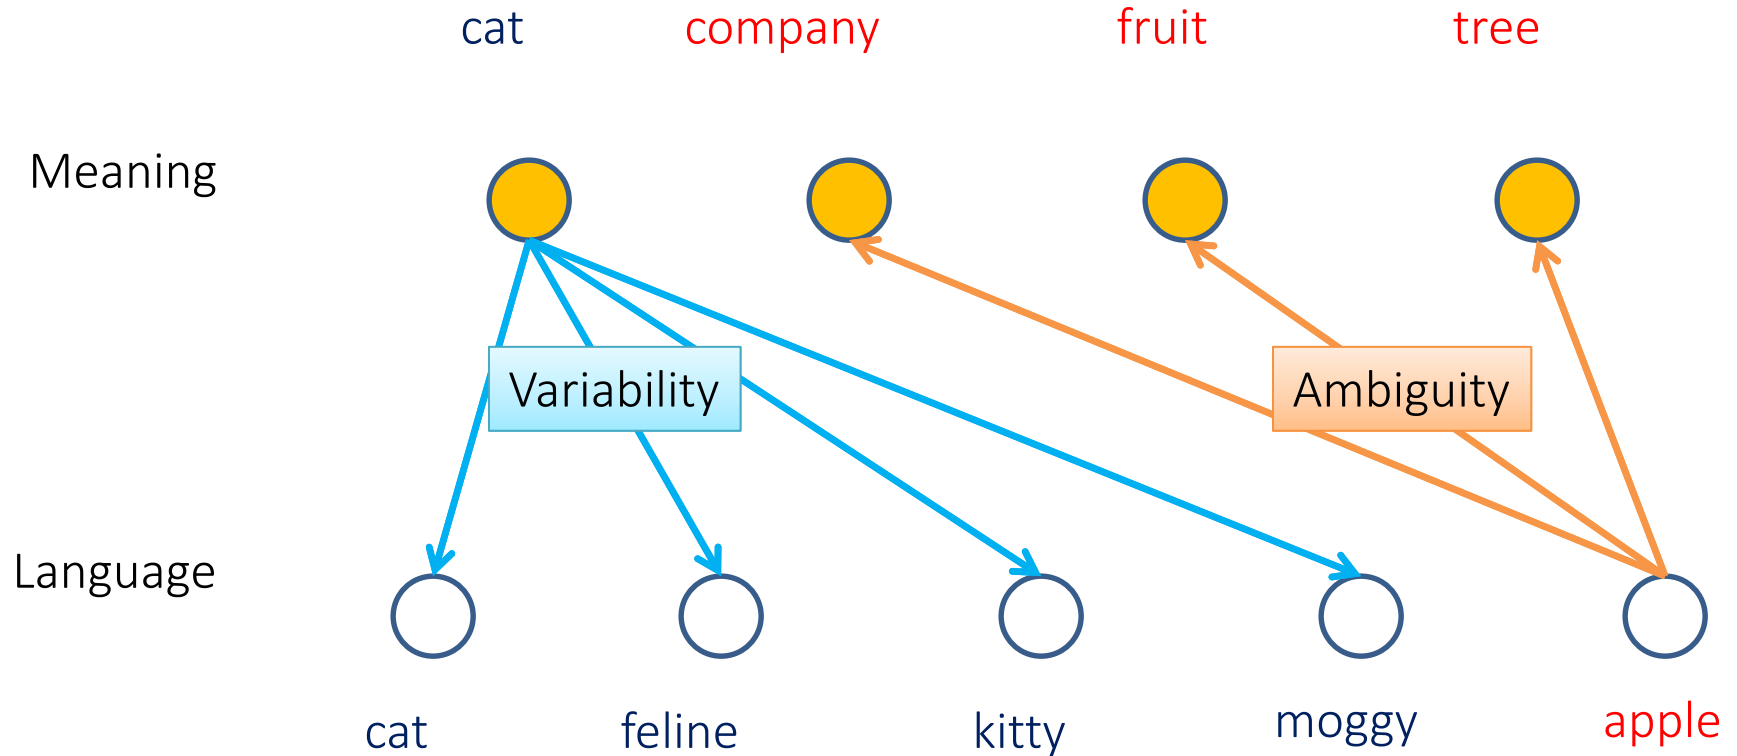


Toy



Why is NLP Difficult?

- Variability and ambiguity everywhere



Words are Ambiguous (have multiple meanings)

- I know that.
- I know that block.
- I know that blocks the sun.
- I know that block blocks the sun.

More Examples of Ambiguity

- Get the cat with the gloves.



Language Subtleties

- Adjective order and placement

- A big black dog
- A big black scary dog
- A big scary dog
- A scary big dog
- ✗ A black big dog

- Antonyms

- Which sizes go together?
 - Big and little
 - Big and small
 - Large and small
 - ✗ Large and little

Levels of Linguistic Analysis: Analogy with Programming Languages

Pragmatics: what does it do?

← implemented the right algorithm

↑
Semantics: what does it mean?

← no implementation bugs

↑
Syntax: what is grammatical?

← no compiler errors

↑
Morphology: basic unit of words

← naming your world

Analogy with Programming Languages

- **Syntax**: no compiler errors
- **Semantics**: no implementation bugs
- **Pragmatics**: implemented the right algorithm
- Different **syntax**, same **semantics** (5):
 - $2 + 3 \leftrightarrow 3 + 2$
- Good **semantics**, bad **pragmatics**:
 - correct implementation of deep neural network
 - for estimating coin flip prob.


How to do natural language processing?

The Role of Memorization

- Children learn words quickly
 - As many as 9 words/day
 - Often only need one exposure to associate meaning with word
 - Can make mistakes, e.g., overgeneralization
 - “I goed to the store.”
 - Exactly how they do this is still under study

The Role of Memorization

- Dogs can do word association too!
 - Rico, a border collie in Germany
 - Knows the names of each of 100 toys
 - Can retrieve items called out to him with over 90% accuracy.
 - Can also learn and remember the names of unfamiliar toys after just one encounter, putting him on a par with a three-year-old child.



[https://en.wikipedia.org/wiki/Rico_\(dog\)](https://en.wikipedia.org/wiki/Rico_(dog))

But there is too much to memorize!

establish

establishment

the church of England as the official state church.

disestablishment

antidisestablishment

antidisestablishmentarian

antidisestablishmentarianism

is a political philosophy that is opposed to the separation of church and state.

Rules and Memorization

- Current thinking in psycholinguistics is that we use a combination of rules and memorization
 - However, this is very controversial
- Mechanism:
 - If there is an applicable rule, apply it
 - However, if there is a memorized version, that takes precedence. (Important for irregular words.)
 - Artists paint “still lifes”
 - Not “still lives”
 - Past tense of
 - think → thought
 - blink → blinked

Representation of Meaning

- I know that block blocks the sun.
 - How do we represent the meanings of “block”?
 - How do we represent “I know”?
 - How does that differ from “I know that.”?
 - Who is “I”?
 - How do we indicate that we are talking about earth’s sun vs. some other planet’s sun?
 - When did this take place? What if I move the block? What if I move my viewpoint? How do we represent this?

How to tackle these problems?

- The field was stuck for quite some time.
- A new approach started around 1990
 - Well, not really new, but the first time around, in the 50's, they didn't have the text, disk space, or GHz
- Main idea: combine memorizing and rules
- How to do it:
 - Get large text collections (corpora)
 - Compute statistics over the words in those collections
- Surprisingly effective
 - Even better now with the Web

NLP ?= Machine Learning

- To be successful, a machine learner needs bias/assumptions; for NLP, that might be linguistic theory/representations.
- Computer representation of language is not directly observable.
- Early connections to information theory (1940s)
- Symbolic, probabilistic, and connectionist ML have all seen NLP as a source of inspiring applications.

NLP \neq Linguistics

- NLP must contend with NL data as found in the world
- NLP \approx computational linguistics
- Linguistics has begun to use tools originating in NLP!

Desiderata for NLP Methods

(ordered arbitrarily)

- Sensitivity to a wide range of the phenomena and constraints in human language
- Generality across different languages, genres, styles, and modalities
- Computational efficiency at construction time and runtime
- Strong formal guarantees (e.g., convergence, statistical efficiency, consistency, robustness, etc.)
- High accuracy when judged against expert annotations and/or task-specific performance

Fields with Connections to NLP

- Machine learning
- Linguistics (including psycho-, socio-, descriptive, and theoretical)
- Cognitive science
- Information theory
- Logic
- Theory of computation
- Data science
- Social and political science
- Psychology
- Economics
- Education

Fields with Connections to Machine Learning

- NLP
- Data mining/Data science
- Bioinformatics
- Fintech
- Computer vision
- Multimedia analysis
- Social and political science
- Psychology
- Economics
- Education


This Course: Machine learning for NLP

- mid-1970s: **HMMs** for speech recognition → probabilistic models
- early 2000s: **conditional random fields** for part-of-speech tagging → structured prediction
- early 2000s: **latent Dirichlet allocation** for modeling text documents → topic modeling
- mid 2010s: **sequence-to-sequence** models for machine translation → neural networks with memory/state


We will select some of the important topics

Nowadays: Deep learning for NLP


- Sequence models



- Memory models



- Attention models



Real-World Applications of NLP

- Spelling suggestions/corrections
- Grammar checking
- Information extraction
- Text categorization
- Automated customer service
 - Conversational agents
 - Question answering
- Speech recognition (limited)
- Machine translation
- Social media analysis
- Rich visual understanding
- Mining legal, medical, or scholarly literature

What We'll Do in this Course

- Learn fundamental machine learning models for NLP
 - Classification models
 - Language models
 - Sequence labeling models
 - Advance NLP tasks using deep learning
- Use NLTK (Natural Language ToolKit) and Tensorflow/PyTorch to try out various algorithms
 - Some assignment will be to do some exercises

What We'll Do in this Course

- Adopt a large text collection
- Use a wide range of NLP techniques to process it
- Release the results for others to use

How to analyze a big collection?

- Your ideas go here with a project

Labs

- No lab in the first week
- TAs will give you some illustrations to show you how to implement algorithms dealing with NLP problem
- You will be asked to submit your results after the tutorial as assignments.
- Some free computational resources to use
 - <https://www.kaggle.com/docs/notebooks>
 - <https://www.kaggle.com/dansbecker/running-kaggle-kernels-with-a-gpu>

Course Information

- Work load and grading:
 - Assignments (30%)
 - Mostly for lab, sometimes a quiz released in Canvas
 - Very simple, you will get points by trying our code and submitting results
 - Projects and Presentation (30%)
 - Project (20%): a team based project
 - For students enrolled in Math4824B, [you can select to do a survey paper](#).
 - Proposal (5%): the proposal deadline will be around mid-term
 - Up to **3 students** in each group
 - In-class Presentation (10%): present what you have done in the project/survey
 - ~~– Midterm~~
 - Final exam (40%)
 - Academic integrity policy
 - Late submission: score got reduced (time based)
 - Plagiarism: all involved parties will get zero

Implementation of Bonus Points

- For each student, up to **three bonus credit tickets** can be earned by
 - Step 1: Answer a question in the class
 - Step 2: After the class, you need to **send me an email** and let me know your name and your name shown in Zoom which I can check the recorded video
- Application of the tickets
 - Add 1 out of 100 points
 - Maximally you can get 3 points to boot you with one level
 - E.g., A → A+